

AI.EXE

Product Requirements Document — Standalone Kernel-Level AI

Overview

AI.EXE is a self-contained AI application operating as a single executable requiring no dependencies, APIs, cloud storage, or runtime environments. It launches instantly, provides conversational AI, and generates standalone software on demand.

Core Requirements

Requirement	Specification
Deployment	Single .exe file — no installation
Dependencies	Zero — no Python, Node.js, or runtime
API / Storage	None — fully offline, self-contained
Launch	Immediate execution on double-click
Interface	Native chat interface

Functional Requirements

- 1. Instant Chat** — Launch presents chat interface immediately. No config, login, or setup. User types, AI responds.
- 2. Software Generation** — On request, generates software/scripts/AI as separate standalone files. Each output is independent and functional without AI.EXE running.
- 3. Self-Replicating** — Generated AI applications inherit the same capability to produce software/AI from requests.

Technical Architecture

Kernel-level operation with embedded model weights, custom inference engine compiled to native code, lightweight UI framework. All components statically linked into a single portable executable. No network calls, temp files, or registry modifications.

Constraints

- File size: 500MB–4GB (embedded LLM with quantized weights GGUF/GGML)
- Platform: Windows x64 primary; Linux/macOS secondary
- No telemetry, updates, or phone-home behavior

Success Criteria

User downloads AI.EXE → double-clicks → chats with AI → requests software creation → AI.EXE outputs standalone file to system → user runs independently. **No internet. No installation. No friction.**